

Model selection



Overview

Quick review of multicollinearity and polynomial regression

Model selection

- Overfitting
- Statistics and methods useful for model selection

If there is time

- Using ggplot to visualize regression models
- Regularization methods

Review: Multicollinearity

Multicollinearity occurs when two or more variables are closely related to each other

- e.g., if they have a high correlation


Multicollinearity can make our estimate of the regression coefficients unstable

- e.g., standard error of coefficients become large


The **variance inflated factor** is a statistic that can be computed to test for multicollinearity

- Rule of thumb: suspect multicollinearity for $VIF > 5$

Left foot



Right foot



`car::vif(lm_fit)`

Review: Polynomial regression

Polynomial regression extends linear regression to non-linear relationships by including nonlinear transformations of predictors

$$\begin{aligned}\text{salary} = & \beta_0 + \beta_1 \cdot \text{endowment} \\ & + \beta_2 \cdot (\text{endowment})^2 + \\ & + \beta_3 \cdot (\text{endowment})^3 + \varepsilon\end{aligned}$$

Still a linear equation but non-linear in original predictors

Review: Polynomial regression



Polynomial regression extends linear regression to non-linear relationships by including nonlinear transformations of predictors

We can compare model fits by:

- Assessing if the coefficients on the higher order terms are statistically significant
- Looking at the R^2 values
- Running hypothesis tests comparing nested models
- Etc.

Let's review this in R...

PREDICT TRAFFIC JAMS



REGRESSION

Model selection

Model selection

Model selection is the process of selecting a statistical model from a set of candidate models

- E.g., which explanatory variables, interaction terms, transformations of variables, etc. to include in a final model

Model selection is a bit of an art

- “All models are wrong but some are useful”




Model selection

Model selection depends on our goal, which usually is either:

- Making accurate predictions
- Understanding relationships in our data


If we fit our model too closely to the data sample we have, we are likely to fail in both of these goals

- i.e., it will be hard to understand relationships between explanatory and response variables
- and model will not make good predictions on new data



Overfitting



Overfitting occurs when we generate a function that too closely matches random sample we have, but does not generalize to the full probability distribution



Overfitting

Overfitting occurs when we generate a function that too closely matches random sample we have, but does not generalize to the full probability distribution

- The model is fit too closely to the shadows and not getting at the Truth



Overfitting song



<https://www.youtube.com/watch?v=DQWI1kvmwRg>

Selecting models methods

There are a number of different methods for selecting models. Four we will briefly discuss are:

1. Creating measures of fit (statistics) that penalize models with more predictors
2. Creating simpler models by removing predictors
3. Evaluating models using cross-validation
4. If there is time: methods that shrink regression coefficients

Model method selection 1


Selected models using statistics that penalize larger models

R^2 as a measure of model fit

We have used the coefficient of multiple determination (R^2) to determine how well our model is fitting the data:

$$R^2 = \frac{SS_{Model}}{SS_{Total}} = 1 - \frac{SS_{Residual}}{SS_{Total}} = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y})^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$$

R^2 always increases with more predictors x_i because the response variable y can always fit more closely with more predictors




Recall: the standard deviation of the errors: σ_ε

Recall for simple linear regression, the standard deviation of the errors is denoted σ_ε and shows how far the points fall off the true regression line.

We can use the **standard deviation of residuals** ($\hat{\sigma}_e$) as an estimate for σ_ε

For simple linear regression we had:


$$\begin{aligned}\hat{\sigma}_e &= \sqrt{\frac{1}{n-2} SSRes} \\ &= \sqrt{\frac{1}{n-2} \sum_{i=1}^n (y_i - \hat{y}_i)^2}\end{aligned}$$



Recall: the standard deviation of the errors: σ_ε

For multiple regression, we with k parameters (i.e., $k - 1$ predictors) an (almost) unbiased estimate of $\hat{\sigma}_\varepsilon$ is:

$$\begin{aligned}\hat{\sigma}_e &= \sqrt{\frac{1}{n-k} SSRes} \\ &= \sqrt{\frac{1}{n-k} \sum_{i=1}^n (y_i - \hat{y}_i)^2}\end{aligned}$$



The residual standard deviation $\hat{\sigma}_\varepsilon$ corrects for bias by dividing the SSResidual by $1/(n - k)$
This estimate does not always decrease with more predictors x

Adjusted R^2

The **adjusted R^2** helps account for the number of predictors in the model by using $\hat{\sigma}_\epsilon^2$

$$R_{adj}^2 = 1 - \frac{SSRes / (n - k)}{SSTotal / n - 1} = 1 - \frac{\hat{\sigma}_e^2}{s_y^2}$$

The adjusted R^2 does not always give a higher values to the model with the more predictors

- i.e., using this statistic, we will not always say that a model with the most predictors is a “better” fit to the data

Other statistics that penalize larger models

There are several other statistics that also ***penalize models that have more predictors***

- These statistics are only meaningful for within data set comparisons

Akaike information criterion: $AIC \propto 2 \cdot k + n \cdot \ln(SSRes/n)$ R: `AIC(lm_fit)`

Bayesian information criterion: $BIC \propto k \cdot \ln(n) + n \cdot \ln(SSRes/n)$ R: `BIC(lm_fit)`

One should select the model with the lowest value on these statistics

Let's try it in R...

Model method selection 2

Using algorithms to select a subset of variables
(stepwise regression)

Brief mention: Variable selection

Variable selection refers to finding models that rely on a small subset of predictors

- This can help make the regression model more interpretable as well

We could use individual predictor's (x's) p-values to determine which predictors to use, however...

- Some of these will be spuriously significant
 - i.e., if H_0 is true for all predictors, ~5 will be significant at $\alpha = .05$ level
- The p-values change as predictors are added and removed
 - Due to multicollinearity

```
lm_fit_mult <-  
  lm(log(endowment) ~  
    salary_tot,  
    salary_men,  
    salary_women  
  )
```

Feature selection: deciding which variables to use

Ideally, we would like to try all combinations of predictors, however, if there are k features, there are 2^k possible models which can be intractable

A few heuristic methods exist for selecting smaller models

- Forward selection: start with a model with no predictors and add predictors (until you have enough)
- Backward selection: Start with the full model and delete predictors
- Mixed selection: Use a combination of forward and backward selection

Coefficients:




	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	12.85948900	0.17913330	71.787	< 0.00000000000000002 ***
salary_tot	0.00013686	0.00003950	3.465	0.000551 ***
salary_men	-0.00001569	0.00002019	-0.777	0.437160
salary_women	-0.00004304	0.00002156	-1.997	0.046099 *

In R: `leaps::regsubsets()`



Model method selection 3

Choosing a model through cross-validation

Cross-validation



Training set



Test set



Good predictions on the test set give an estimate of how accurate the model will be on new data from the population

Cross-validation

We run cross-validation by splitting data into two sets:

A training set in which the parameters of a regression model are fit (estimated)

A test set in which the prediction accuracy of our model is assessed



Mean squared prediction error

To evaluate how effective a model is, we can use the mean squared prediction error (MSPE) using the following steps:

1. Fit a model using the training data
2. Make predictions on the test data
3. We can use the MSPE to assess how accurate the predictions are:

$$MSPE = \frac{1}{n_t} \sum_{i=1}^{n_t} (y_i - \hat{y}_i)^2 = \frac{1}{n_t} \sum_{i=1}^{n_t} (y_i - (\hat{\beta}_0 + \hat{\beta}_1 x_1 + \dots + \hat{\beta}_k x_k))^2$$

Actual y values in the **test set**

Predicted y values on the **test set**

Parameters estimated on the training set


Predictions assessed on the test set

n_t is the number of points in the test set

K-fold cross-validation

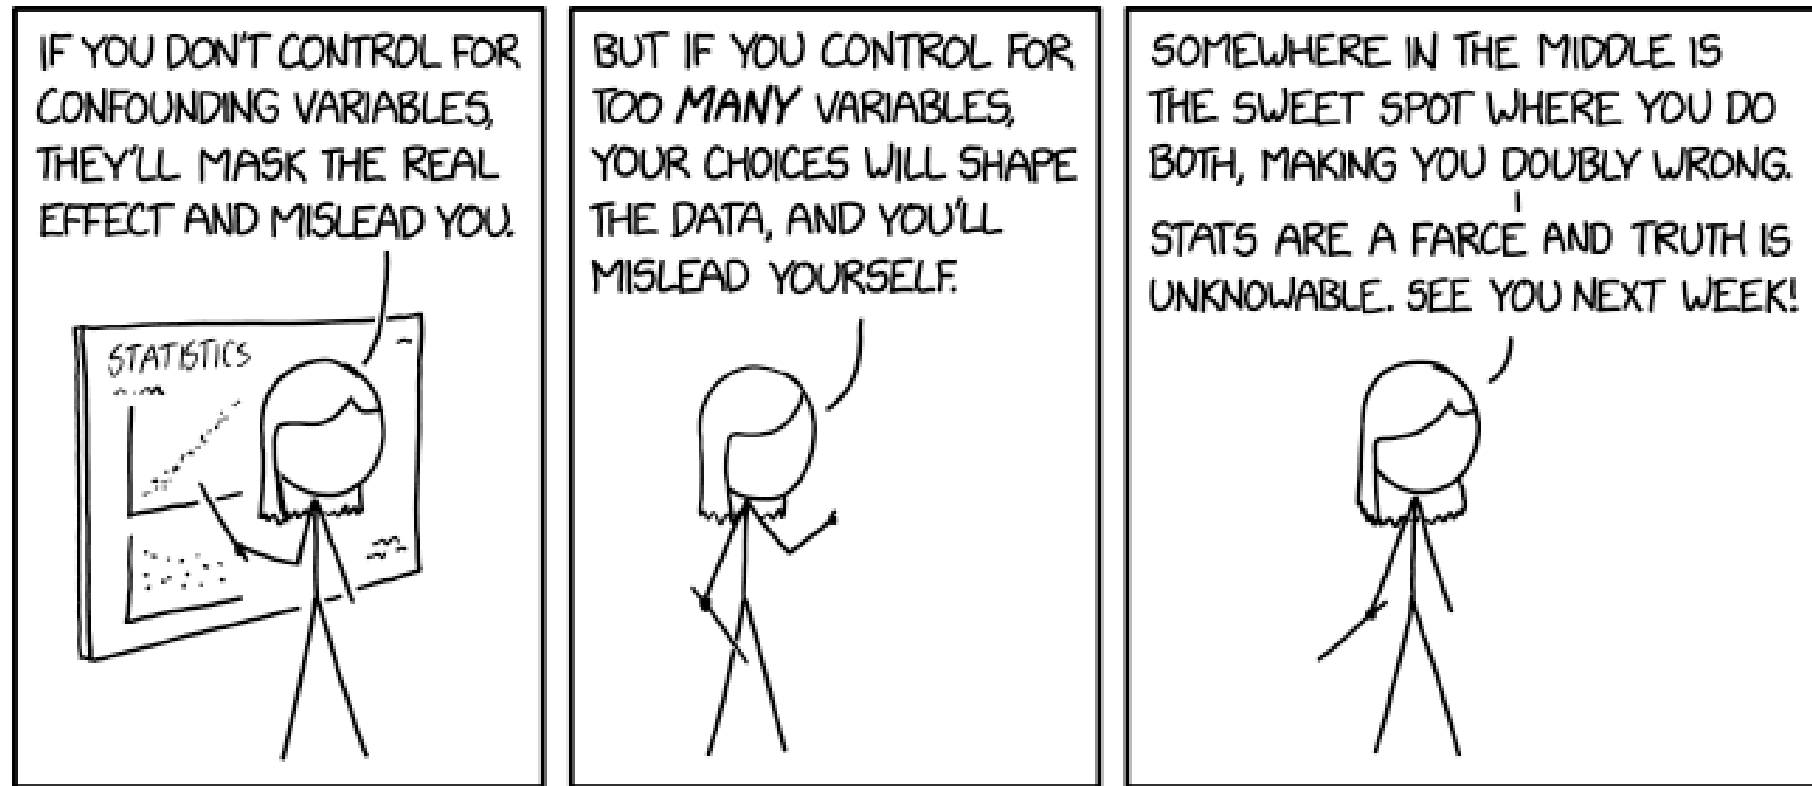
K-fold cross-validation

- Split the data into k parts
- Train on $k-1$ of these parts and test on the left out part
- Repeat this process for all k parts
- Average the prediction accuracies to get a final estimate of the generalization error



Leave-one-out (LOO)
cross-validation: $k = n$

Let's try it in R...




Plotting multiple regression models with ggplot

So far we have plotted our multiple regression models using base R graphics

This was useful for seeing the relationship between how R fits linear models, and what these models represent

However, if you want an easier/prettier way to visualize linear models, we can use ggplot!



```
Call:
lm(formula = salary_tot ~ log_endowment + rank_name + log_endowment:rank_name,
    data = IPED_2)
```

Let's try it in R!

```
Residual standard error: 16260 on 705 degrees of freedom
Multiple R-squared: 0.7806, Adjusted R-squared: 0.7791
F-statistic: 501.7 on 5 and 705 DF, p-value: < 2.2e-16
```

Side note

Controlling model complexity with regularization

Brief mention: shrinkage methods

Rather than finding the coefficients $\hat{\beta}_i$ by just minimizing the SSRes, one can also add penalties in the fitting procedure to find simpler models



We will very briefly discuss two techniques:

- Ridge regression (l_2 norm penalty)
- The lasso (l_1 norm penalty)

Brief mention: Ridge regression

Ridge regression finds the coefficients $\hat{\beta}_i$ that minimize:

$$\sum_{i=1}^n (y_i - (\underbrace{\hat{\beta}_0 + \sum_{j=1}^k \hat{\beta}_j x_{ij}}_{\text{SSRes}}))^2 + \lambda \underbrace{\sum_{j=1}^k \hat{\beta}_j^2}_{\text{shrinkage penalty}}$$

Tuning parameter

What happens if:

- $\lambda = 0$
- $\lambda \rightarrow \text{infinity}$
- (the coefficients depend on the tuning parameter value)


see the glmnet package

Brief mention: The Lasso

The Lasso finds the coefficients $\hat{\beta}_i$ that minimize:

$$\underbrace{\sum_{i=1}^n (y_i - (\hat{\beta}_0 + \sum_{j=1}^k \hat{\beta}_j x_{ij}))^2}_{\text{SSRes}} + \lambda \underbrace{\sum_{j=1}^k |\hat{\beta}_j|}_{\text{shrinkage penalty}}$$

Tuning parameter



Similar to ridge regression but penalizes $|\hat{\beta}_i|$ instead of $\hat{\beta}_i^2$

- i.e., uses the L_1 penalty instead of the L_2 penalty

Advantages

- Final model will often have many set $\hat{\beta}_i$ to 0
- i.e., does variable selection and creates a 'sparse' model

see the glmnet package

Shrinkage/regularization methods

Regularization methods often work very well when we care about making accurate predictions on new data

Theory suggests that these methods work by minimizing the MSPE through a bias-variance tradeoff

- Average MSPE = $\text{bias}^2 + \text{variance}$
- We use a biased method (via regularization) to get models that vary less from one random data set to the next, reducing the average MSPE

To learn more about regularization methods, take a class on Machine Learning!